

QIUSHI LYU

Peking University

@lvqiushi2015@gmail.com, lvqiushi@stu.pku.edu.cn
447, Building 45B, Peking University, Beijing, 100871

(+86)136 7518 5585
Beijing, China

<https://alfredlyu.github.io/>

AlfredLyu

EDUCATION

Bachelors of Computer Science and Technology (Turing Class)

Peking University

Sept. 2021 – Present

Beijing, China

- GPA 3.92/4.00, **Rank 1/121**
- Gold medalist at the 2020 China National Olympiad in Informatics (NOI), leading to direct admission into Peking University.
- Entered Peking University for preparatory studies from Mar. 2021 to Aug. 2021.

Visiting Student

Massachusetts Institute of Technology

Mar. 2024 – Aug. 2024

Boston, US

- Advised by Prof. Joshua B. Tenenbaum and Prof. Chuang Gan.
- Specialized in Computer Vision and Embodied AI.

PUBLICATIONS (*DENOTES EQUAL CONTRIBUTION)

Constrained Human-AI Cooperation: An Inclusive Embodied Social Intelligence Challenge [link]

NeurIPS 2024 D & B Track

Authors: Weihua Du*, **Qiushi Lyu***, Jiaming Shan, Zhenting Qi, Hongxin Zhang, Sunli Chen, Andi Peng, Tianmin Shu, Kwongjoon Lee, Behzad Dariush, Chuang Gan

- We created a new benchmark, named CHAIC, to test embodied agents' ability to actively perceive human partners' intents and constraints from egocentric visual observations. We designed new agents with real physical constraints and some long-horizon tasks featuring both indoor and outdoor scenes.

COMBO: Compositional World Models for Embodied Multi-Agent Cooperation [link]

Submit to ICLR 2025

Authors: Hongxin Zhang*, Zeyuan Wang*, **Qiushi Lyu***, Zheyuan Zhang, Sunli Chen, Tianmin Shu, Yilun Du, Chuang Gan

- ICLR 2025 average review score 6.67 (8, 6, 6)
- We introduced an embodied multi-agent planning framework that leverages a compositional world model to empower the agents to imagine how different actions may affect the world in the long run and plan more cooperatively. The compositional world model is learned by factorizing joint actions of agents and compositionally generating the future frames of the world state.

2.5D video-audio generation from single image

In Progress

Authors: Sunli Chen*, **Qiushi Lyu***, Haotian Yuan*, Kaizhi Qian, Yang Zhang, Chuang Gan

- We created a pipeline to generate stereo sound from silent videos and used it to generate a high-quality video dataset with stereo sound. We trained a model on this dataset that jointly generates video and stereo sound given an initial frame.

MagnifierSketch: A Unified Data Structure for Efficient Aggregate and Per-Flow Latency Quantile Estimation

Submit to WWW 2025

Authors: Jiarui Guo*, **Qiushi Lyu***, Yuhan Wu*, Haoyu Li*, Tong Yang, Zhaoqian Yao, Yuqi Dong, Peiqing Chen, Xiaolin Wang

- We proposed MagnifierSketch, an efficient algorithm for per-flow latency quantile estimation, which is unbiased and significantly better than the state-of-the-art algorithms.

RESEARCH INTERESTS

Computer Vision

Embodied AI and Robotics

Multimodal Learning






Cognitive Science

LANGUAGES

Chinese: Native Speaker

English: TOEFL 105(Speaking 23), GRE 320

SELECTED PRIZES

	May 4th Scholarship The best scholarship in Peking University, has equivalent status to National Scholarship	Dec. 2023
	POSCO Asia Fellowship Comprehensive Excellent Award sponsored by POSCO	Dec. 2022
	International Collegiate Programming Contest (ICPC) Asia-East Continent Final Contest Rank 4, Gold Medal	Jul. 2022
	International Collegiate Programming Contest (ICPC) Shenyang Regional Contest Rank 1 , Gold Medal, beat the ICPC World Final Champion	Nov. 2021
	China National Olympiad in Informatics (NOI) 2020 Rank 45, Gold Medal	Aug. 2020

EXPERIENCES

Research Intern

MIT-IBM Watson AI Lab, Massachusetts Institute of Technology

📅 Aug. 2023 – Present

📍 Boston, US / Remote

- Advised by Prof. Joshua B. Tenenbaum and Prof. Chuang Gan.
- Developed *COMBO*, an embodied multi-agent planning framework that leverages a compositional world model to empower the agents to imagine future actions and plan more cooperatively.
- Also Developed *CHAIC*, a new benchmark to test embodied agents' ability to actively perceive human partners' intents and constraints from egocentric visual observations and cooperate with them more efficiently.
- Another work aimed to generate a high-quality video dataset with stereo sound and train a model on this dataset that jointly generates video and stereo sound given an initial frame.

Research Intern

Peking University

📅 Jan. 2023 – July. 2023

📍 Beijing, China

- Advised by Prof. Tong Yang
- Focused on sketch algorithms, introduced *MagnifierSketch*, an efficient algorithm for per-flow latency quantile estimation.

TEACHING EXPERIENCE

Teaching Assistant: Discrete Mathematics and Structures(I)

Fall 2023